

frozen_lake

November 18, 2023

1 Frozen Lake Assignment

Grzegorz Malisz Student Number: 4852370

1.1 Imports

```
[12]: import gymnasium as gym
import numpy as np
import matplotlib.pyplot as plt

MOVING = True # extra cleaner plot with moving-average of reward
if MOVING:
    from scipy.signal import savgol_filter
```

1.2 Sample code as function

```
[13]: def reinforced_learning_and_plots(learning_rate=0.1,
                                         discount_factor=0.9,
                                         episodes=4 * 1000,
                                         t=100,
                                         exploration_rate=0.01,
                                         cost_of_living=- 0.01,
                                         frequent_rewards=False,
                                         is_slippery=False):

    # # hyper parameters
    # learning_rate = 0.1 #impacts how fast we update our estimates
    # discount_factor = 0.9 # gamma, impacts the return calculations
    # episodes = 4*1000 # number of episodes to learn, keep this a multiple of 4
    # four for nice plotting
    # t = 100 # Maximum steps in an episode
    # exploration_rate = 0.5 # Exploration rate
    # cost_of_living = - 0.01 # used when frequent_rewards = True, incentive
    # the agent for efficiency by incurring a cost to each move
    # frequent_rewards = False # when False the original environment rewards
    # are used, if True the agent is incentivised to be efficient and avoid holes
    # (how?)
    # is_slippery = False
```

```

# Choose environment

env = gym.make("FrozenLake8x8-v1", is_slippery=is_slippery)

Q = np.zeros([env.observation_space.n, env.action_space.n])
rewards_per_episode = []
q_values_at_intervals = [] # Store Q-values at intervals

for episode in range(epochs):
    state, prob = env.reset()
    total_reward = 0

    for step in range(t):
        # Choose action based on epsilon-greedy policy
        if np.random.rand() < exploration_rate:
            action = np.random.choice(env.action_space.n)
        else:
            action = np.argmax(Q[state, :])

        new_state, reward, terminated, truncated, info = env.step(action)

        if frequent_rewards:
            if terminated & (reward == 0):
                reward = reward - 1

            reward = reward + cost_of_living

        # Update Q-value using Q-learning equation
        Q[state, action] += learning_rate * (reward + discount_factor * np.
↪ max(Q[new_state, :]) - Q[state, action])
        total_reward += reward
        state = new_state

        if terminated:
            #if reward == 1 + cost_of_living:
            # print(f"Episode {episode} finished after {step + 1} steps.
↪ Success!")
            break

    rewards_per_episode.append(total_reward)

    # Store Q-values at intervals (e.g., every 100 episodes)
    if (episode + 1) % (epochs // 4) == 0:
        q_values_at_intervals.append(np.copy(Q)) # Store a copy of Q-values

# Plotting rewards per episode
plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)

```

```

plt.plot(rewards_per_episode, label='Total Reward')
plt.title('Rewards per Episode')
plt.xlabel('Episode')
plt.ylabel('Total Reward')

# the plot is too noisy, you can use the scipy package to calculate a
↪moving average
if MOVING:
    moving_average_window = 20
    moving_averages = savgol_filter(rewards_per_episode,
↪moving_average_window, 3)
    plt.plot(moving_averages, label=f'Moving Average (Window
↪{moving_average_window})', color='orange')

plt.legend()

# Plotting the heatmap of Q-values at intervals
fig, ax = plt.subplots(1, len(q_values_at_intervals), figsize=(15, 5))

for i, q_values in enumerate(q_values_at_intervals):
    ax[i].imshow(q_values, cmap='hot', interpolation='nearest')
    ax[i].set_title(f'Q-Values at Episode {(episodes // 4) * (i + 1)}')
    ax[i].axis('off') # Turn off axis
    plt.colorbar(ax[i].imshow(q_values, cmap='hot',
↪interpolation='nearest'), ax=ax[i])
    plt.pause(0.1) # Pause briefly to update the plot
    break

```


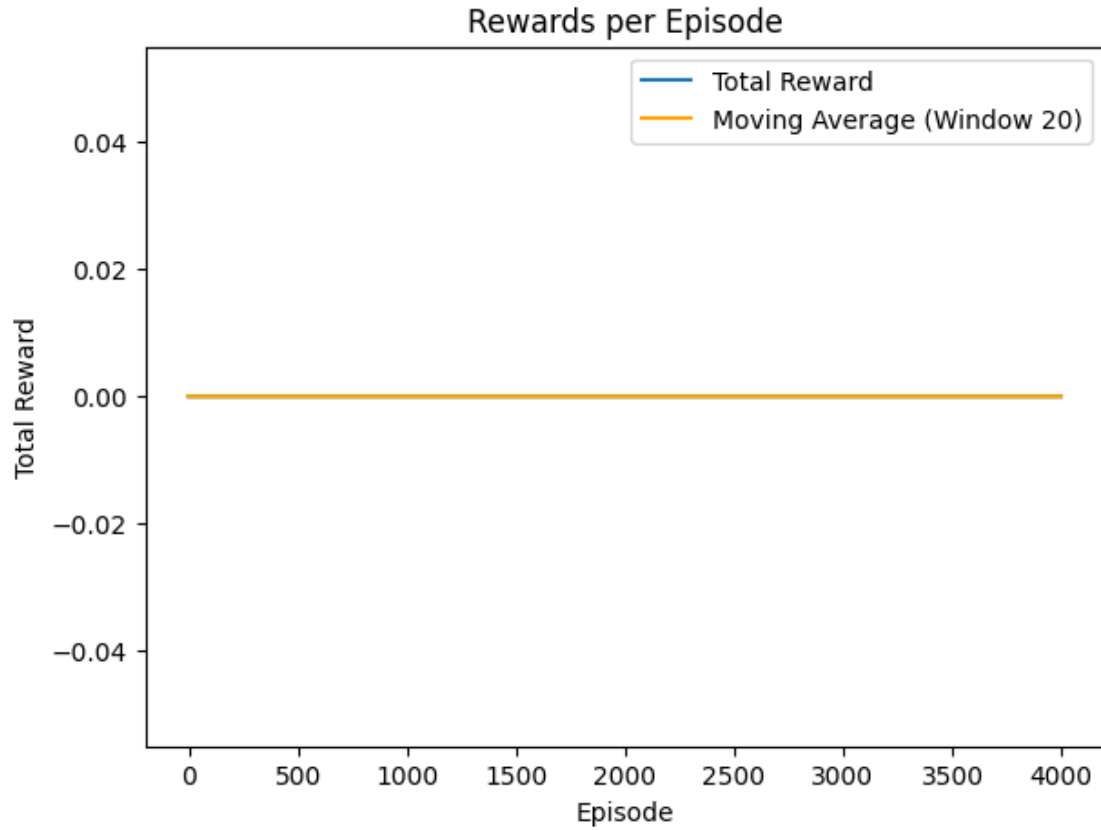
1.3 t=100, exploration_rate=0.01, frequent_rewards=False,
is_slippery=False

This makes the agent do mostly action from the 0 index, and agent does not learn anything.

```

[14]: reinforced_learning_and_plots(t=100,
                                     exploration_rate=0.01,
                                     frequent_rewards=False,
                                     is_slippery=False)

```




1.4 `t=100,` `exploration_rate=0.01,` `frequent_rewards=True,`
`is_slippery=False`

This makes the agent change his tactic, because he is penalised for wrong moves (`exploration_rate` is low) he is trying to get the path with max reward. Frequent Learning had huge impact because

it allowed agent to finally start updateing q-table.

```
[15]: reinforced_learning_and_plots(t=100,  
                                     exploration_rate=0.01,  
                                     frequent_rewards=True,  
                                     is_slippery=False)
```




1.5 `t=1000,` `exploration_rate=0.01,` `frequent_rewards=True,`
`is_slippery=False`

Changing T to 1000, did not change much agent behaviour, only his first moves have lower reward.

```
[16]: reinforced_learning_and_plots(t=1000,  
                                   exploration_rate=0.01,  
                                   frequent_rewards=True,  
                                   is_slippery=False)
```







1.6 `t=100, exploration_rate=0.01, frequent_rewards=True, is_slippery=True`

Huge reward noise, but still agent is learning.


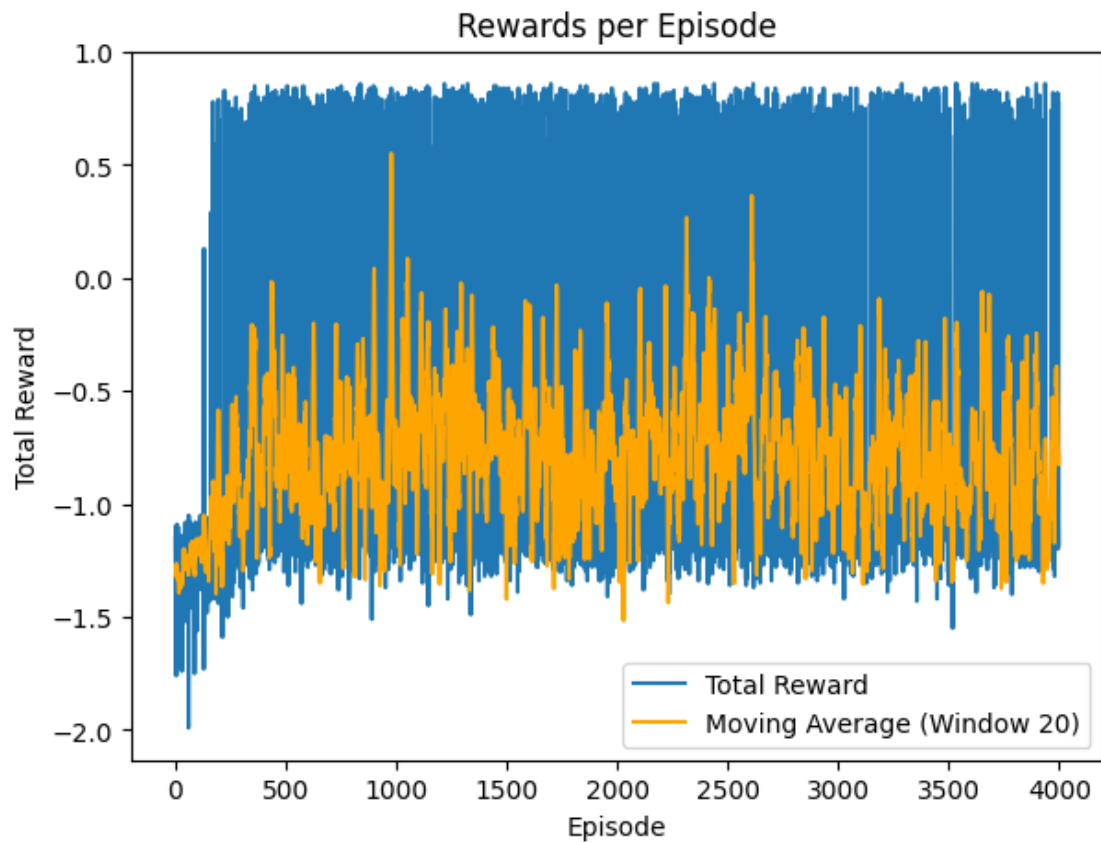
```
[17]: reinforced_learning_and_plots(t=100,  
                                     exploration_rate=0.01,  
                                     frequent_rewards=True,  
                                     is_slippery=True)
```



1.7 `t=100, exploration_rate=0.5, frequent_rewards=True, is_slippery=False`

Moving average reward is always 0, due to agent choosing practically randomly. The fact that Q-learning is optimistic means that it is choosing a path with the highest potential reward.


```
[18]: reinforced_learning_and_plots(t=100,
                                     exploration_rate=0.5,
                                     frequent_rewards=True,
                                     is_slippery=False)
```



1.8 Q-learning over thoughts

Using each experience only once to update the Q-table is disadvantage for rare events, because it leads to suboptimal or inaccurate Q-values (not taking advantage/avoiding of rare events). It can be a problem especially with small amount of learning episodes or small exploration rate.